

```

import math
import os

import numpy as np
import pandas as pd
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline
from imblearn.under_sampling import RandomUnderSampler
from sklearn import preprocessing

from sklearn.compose import ColumnTransformer, make_column_selector
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MultiLabelBinarizer

os.environ['LOKY_MAX_CPU_COUNT'] = str(os.cpu_count()-1) # To silence
warning : Could not find the number of physical cores

def allOutliersToBound(data):
    dfOutput = data.copy()

    for col in dfOutput.columns:
        if dfOutput[col].dtype == 'int64': # The columns with 'int64'
            come from encoded data and only having values of 0 and 1, function may
            cause every data become same value
            continue
        if col == 'imdb_rating':           # Prevent adjusting the
target
            continue
        outliersToBound(dfOutput, col)

    return dfOutput

def outliersToBound(data, col):
    # Calculate 25% and 75% quarter, IQR and bound values
    Q1 = data[col].quantile(0.25)
    Q3 = data[col].quantile(0.75)
    IQR = Q3 - Q1
    lowerBound = Q1 - 1.5 * IQR
    upperBound = Q3 + 1.5 * IQR

    # Cap the outliers to lower bound or upper bound
    data[col] = data[col].clip(lower=lowerBound, upper=upperBound)
    return data

```

Load source file and show data for preview

```

filePath = '../Dataset/movie_metadata.csv'
df = pd.read_csv(filePath)
df

```

	color	director_name	num_critic_for_reviews	duration	\
0	Color	James Cameron	723.0	178.0	
1	Color	Gore Verbinski	302.0	169.0	
2	Color	Sam Mendes	602.0	148.0	
3	Color	Christopher Nolan	813.0	164.0	
4	NaN	Doug Walker	NaN	NaN	
...
5038	Color	Scott Smith	1.0	87.0	
5039	Color	NaN	43.0	43.0	
5040	Color	Benjamin Roberds	13.0	76.0	
5041	Color	Daniel Hsia	14.0	100.0	
5042	Color	Jon Gunn	43.0	90.0	
		director_facebook_likes	actor_3_facebook_likes		
actor_2_name	\				
0	Moore	0.0	855.0	Joel David	
1	Bloom	563.0	1000.0	Orlando	
2	Kinnear	0.0	161.0	Rory	
3	Bale	22000.0	23000.0	Christian	
4	Walker	131.0	NaN	Rob	
...
5038	Zuniga	2.0	318.0	Daphne	
5039	Curry	NaN	319.0	Valorie	
5040	Moody	0.0	0.0	Maxwell	
5041	Henney	0.0	489.0	Daniel	
5042	Herzlinger	16.0	16.0	Brian	
		actor_1_facebook_likes	gross		
genres	\				
0	Sci-Fi	1000.0	760505847.0	Action Adventure Fantasy	
1	Fantasy	40000.0	309404152.0	Action Adventure	
2	Thriller	11000.0	200074175.0	Action Adventure	
3	Thriller	27000.0	448130642.0	Action	
4	Documentary	131.0	NaN		

...
5038	637.0	NaN		Comedy
Drama				
5039	841.0	NaN	Crime Drama Mystery	
Thriller				
5040	0.0	NaN		Drama Horror
Thriller				
5041	946.0	10443.0		Comedy Drama
Romance				
5042	86.0	85222.0		
Documentary				
... num_user_for_reviews language country content_rating				
budget \				
0 ...	3054.0	English	USA	PG-13
237000000.0				
1 ...	1238.0	English	USA	PG-13
300000000.0				
2 ...	994.0	English	UK	PG-13
245000000.0				
3 ...	2701.0	English	USA	PG-13
250000000.0				
4 ...	NaN	NaN	NaN	NaN
NaN				
...
...				
5038 ...	6.0	English	Canada	NaN
NaN				
5039 ...	359.0	English	USA	TV-14
NaN				
5040 ...	3.0	English	USA	NaN
1400.0				
5041 ...	9.0	English	USA	PG-13
NaN				
5042 ...	84.0	English	USA	PG
1100.0				
title_year actor_2_facebook_likes imdb_score aspect_ratio \				
0 2009.0	936.0	7.9	1.78	
1 2007.0	5000.0	7.1	2.35	
2 2015.0	393.0	6.8	2.35	
3 2012.0	23000.0	8.5	2.35	
4 NaN	12.0	7.1	NaN	
...
5038 2013.0	470.0	7.7	NaN	
5039 NaN	593.0	7.5	16.00	
5040 2013.0	0.0	6.3	NaN	
5041 2012.0	719.0	6.3	2.35	
5042 2004.0	23.0	6.6	1.85	

```

    movie_facebook_likes
0                  33000
1                      0
2                  85000
3                 164000
4                      0
...
5038                   84
5039                 32000
5040                   16
5041                   660
5042                   456

```

[5043 rows x 28 columns]

Start data processing

Drop duplicate data

Treat movie_imdb_link as the unique identifier, drop the duplicate rows and check the data again

```

df.sort_values(by='num_voted_users', kind="mergesort", inplace=True)
# Sort by num_voted_users, default ascending
dfNoDuplicate = df.drop_duplicates(subset=['movie_imdb_link'],
keep="last") # Keep the lastest data (assume num_voted_users only
increase along the time)
dfNoDuplicate

```

	color	director_name	num_critic_for_reviews
duration \			
4702	Color	Bill Benenson	1.0
71.0			
4958	Black and White	Harry F. Millarde	1.0
110.0			
279	NaN	Christopher Barnard	NaN
22.0			
4244	Color	Dan Perri	1.0
100.0			
4716	Color	Lance McDaniel	NaN
90.0			
...
...			
3355	Color	Quentin Tarantino	215.0
178.0			
683	Color	David Fincher	315.0
151.0			
97	Color	Christopher Nolan	642.0
148.0			

66	Color	Christopher Nolan	645.0
152.0			
1937	Color	Frank Darabont	199.0
142.0			
director_facebook_likes actor_3_facebook_likes			
actor_2_name \			
4702	0.0	21.0	Dave
Fennoy			
4958	0.0	0.0	Johnnie
Walker			
279	0.0	NaN	
Nan			
4244	0.0	338.0	David
Proval			
4716	0.0	271.0	Steven Michael
Quezada			
...	
...			
3355	16000.0	857.0	Eric
Stoltz			
683	21000.0	637.0	
Meat Loaf			
97	22000.0	23000.0	
Tom Hardy			
66	22000.0	11000.0	Heath
Ledger			
1937	0.0	461.0	Jeffrey
DeMunn			
actor_1_facebook_likes gross \			
4702	1000.0	NaN	
4958	2.0	3000000.0	
279	5.0	NaN	
4244	749.0	NaN	
4716	595.0	NaN	
...	
3355	13000.0	107930000.0	
683	11000.0	37023395.0	
97	29000.0	292568851.0	
66	23000.0	533316061.0	
1937	11000.0	28341469.0	
genres ... num_user_for_reviews			
\			
4702	Documentary	...	1.0
4958	Crime Drama	...	1.0
279	Comedy	...	NaN

4244	Comedy Drama Mystery Romance Thriller	...	1.0			
4716	Action Drama Thriller	...	1.0			
...			
3355	Crime Drama	...	2195.0			
683	Drama	...	2968.0			
97	Action Adventure Sci-Fi Thriller	...	2803.0			
66	Action Crime Drama Thriller	...	4667.0			
1937	Crime Drama	...	4144.0			
	language	country	content_rating	budget	title_year	\
4702	English	USA	NaN	650000.0	2014.0	
4958	NaN	USA	NaN	100000.0	1920.0	
279	NaN	Nan	NaN	NaN	NaN	
4244	English	USA	NaN	2100000.0	2015.0	
4716	English	USA	PG-13	600000.0	2014.0	
...	
3355	English	USA	R	8000000.0	1994.0	
683	English	USA	R	63000000.0	1999.0	
97	English	USA	PG-13	160000000.0	2010.0	
66	English	USA	PG-13	185000000.0	2008.0	
1937	English	USA	R	25000000.0	1994.0	
	actor_2_facebook_likes	imdb_score	aspect_ratio			
movie_facebook_likes						
4702	338.0	7.4	NaN			
5						
4958	2.0	4.8	1.33			
0						
279	NaN	7.2	NaN			
0						
4244	354.0	6.7	2.39			
14						
4716	412.0	8.0	NaN			
9						
...	
...						
3355	902.0	8.9	2.35			
45000						
683	783.0	8.8	2.35			
48000						
97	27000.0	8.8	2.35			

```

175000
66           13000.0      9.0      2.35
37000
1937          745.0      9.3      1.85
108000

```

[4919 rows x 28 columns]

Remove irrelevant and unsuitable columns

Dropping the links (irrelevant) and names (unsuitable)

```

dfValueFeature = dfNoDuplicate.drop(columns=
    ['movie_imdb_link',      # Link is nothing related to the score
     'movie_title',
     'director_name',
     'actor_1_name',
     'actor_2_name',
     'actor_3_name'])       # Title, director and actors might affect
                           # the score, but training a model with 'names' is highly lead to
                           # overfitting
dfValueFeature

```

		color	num_critic_for_reviews	duration	\
4702		Color	1.0	71.0	
4958	Black and White		1.0	110.0	
279		NaN	NaN	22.0	
4244		Color	1.0	100.0	
4716		Color	NaN	90.0	
...		
3355		Color	215.0	178.0	
683		Color	315.0	151.0	
97		Color	642.0	148.0	
66		Color	645.0	152.0	
1937		Color	199.0	142.0	
		director_facebook_likes	actor_3_facebook_likes		
	actor_1_facebook_likes	\			
4702		0.0		21.0	
1000.0					
4958		0.0		0.0	
2.0					
279		0.0		NaN	
5.0					
4244		0.0		338.0	
749.0					
4716		0.0		271.0	
595.0					
...		

...			
3355	16000.0	857.0	
13000.0			
683	21000.0	637.0	
11000.0			
97	22000.0	23000.0	
29000.0			
66	22000.0	11000.0	
23000.0			
1937	0.0	461.0	
11000.0			
gross		genres	
num_voted_users \			
4702	NaN	Documentary	
5			
4958	3000000.0	Crime Drama	
5			
279	NaN	Comedy	
6			
4244	NaN	Comedy Drama Mystery Romance Thriller	
6			
4716	NaN	Action Drama Thriller	
6			
...	
...			
3355	107930000.0	Crime Drama	
1324680			
683	37023395.0	Drama	
1347461			
97	292568851.0	Action Adventure Sci-Fi Thriller	
1468200			
66	533316061.0	Action Crime Drama Thriller	
1676169			
1937	28341469.0	Crime Drama	
1689764			
cast_total_facebook_likes		...	language
country \			
4702	1359	...	1.0 English
USA			
4958	4	...	1.0 NaN
USA			
279	5	...	Nan NaN
Nan			
4244	1814	...	1.0 English
USA			
4716	1754	...	1.0 English
USA			

...
3355		16557	...		2195.0	English
USA						
683		13209	...		2968.0	English
USA						
97		81115	...		2803.0	English
USA						
66		57802	...		4667.0	English
USA						
1937		13495	...		4144.0	English
USA						
	content_rating		budget	title_year	actor_2_facebook_likes	\
4702	NaN	650000.0		2014.0		338.0
4958	NaN	100000.0		1920.0		2.0
279	NaN		NaN			NaN
4244	NaN	2100000.0		2015.0		354.0
4716	PG-13	600000.0		2014.0		412.0
...
3355	R	8000000.0		1994.0		902.0
683	R	63000000.0		1999.0		783.0
97	PG-13	160000000.0		2010.0		27000.0
66	PG-13	185000000.0		2008.0		13000.0
1937	R	25000000.0		1994.0		745.0
	imdb_score	aspect_ratio		movie_facebook_likes		
4702	7.4		NaN		5	
4958	4.8		1.33		0	
279	7.2		NaN		0	
4244	6.7		2.39		14	
4716	8.0		NaN		9	
...	
3355	8.9		2.35		45000	
683	8.8		2.35		48000	
97	8.8		2.35		175000	
66	9.0		2.35		37000	
1937	9.3		1.85		108000	
[4919 rows x 22 columns]						

Fill missing values

Fill 0 or median for numeric data and most frequent value for object data

```
# The minimum value is 1 instead of 0, while the website is showing 0,
# mean most NaN refer to 0
dfValueFeature['num_critic_for_reviews'] =
df['num_critic_for_reviews'].fillna(0)
```

```

# Same as 'num_critic_for_reviews'
dfValueFeature['num_user_for_reviews'] =
df['num_user_for_reviews'].fillna(0)

fillNanTransformer = ColumnTransformer(
    transformers=[
        # Minimum value of 'likes' type data are 0 (NaN might not
        # refer to 0 in these case), gross and budget likely not possible to be
        # 0, fills these columns with median
        ('num', SimpleImputer(strategy='median'),
        make_column_selector(dtype_include=np.number)),
        # Fill by most frequent
        ('cat', SimpleImputer(strategy='most_frequent'),
        make_column_selector(dtype_include=object)),
    ], verbose_feature_names_out=False).set_output(transform="pandas")

```

dfFilledData = fillNanTransformer.fit_transform(dfValueFeature)

Check if any missing value haven't filled

dfFilledData

	num_critic_for_reviews	duration	director_facebook_likes	\
4702	1.0	71.0	0.0	
4958	1.0	110.0	0.0	
279	0.0	22.0	0.0	
4244	1.0	100.0	0.0	
4716	0.0	90.0	0.0	
...	
3355	215.0	178.0	16000.0	
683	315.0	151.0	21000.0	
97	642.0	148.0	22000.0	
66	645.0	152.0	22000.0	
1937	199.0	142.0	0.0	
	actor_3_facebook_likes	actor_1_facebook_likes	gross	\
4702	21.0	1000.0	25035665.0	
4958	0.0	2.0	3000000.0	
279	365.5	5.0	25035665.0	
4244	338.0	749.0	25035665.0	
4716	271.0	595.0	25035665.0	
...	
3355	857.0	13000.0	107930000.0	
683	637.0	11000.0	37023395.0	
97	23000.0	29000.0	292568851.0	
66	11000.0	23000.0	533316061.0	
1937	461.0	11000.0	28341469.0	
	num_voted_users	cast_total_facebook_likes	facenumber_in_poster	\
4702	5.0	1359.0	1.0	

4958	5.0	4.0	1.0
279	6.0	5.0	0.0
4244	6.0	1814.0	0.0
4716	6.0	1754.0	0.0
...
3355	1324680.0	16557.0	1.0
683	1347461.0	13209.0	2.0
97	1468200.0	81115.0	0.0
66	1676169.0	57802.0	0.0
1937	1689764.0	13495.0	0.0
4702	num_user_for_reviews	... actor_2_facebook_likes	imdb_score \
4958	1.0	338.0	7.4
279	0.0	2.0	4.8
4244	1.0	593.0	7.2
4716	1.0	354.0	6.7
...	...	412.0	8.0
3355	2195.0	902.0	8.9
683	2968.0	783.0	8.8
97	2803.0	27000.0	8.8
66	4667.0	13000.0	9.0
1937	4144.0	745.0	9.3
4702	aspect_ratio	movie_facebook_likes	color \
4958	2.35	5.0	Color
279	1.33	0.0	Black and White
4244	2.35	0.0	Color
4716	2.39	14.0	Color
...	...	9.0	Color
3355	2.35
683	2.35	45000.0	Color
97	2.35	48000.0	Color
66	2.35	175000.0	Color
1937	1.85	37000.0	Color
4702	genres \		
4958	Documentary		
	Crime Drama		

```

279                                         Comedy
4244 Comedy|Drama|Mystery|Romance|Thriller
4716             Action|Drama|Thriller
...
3355                 Crime|Drama
683                     Drama
97          Action|Adventure|Sci-Fi|Thriller
66          Action|Crime|Drama|Thriller
1937                 Crime|Drama

plot_keywords language
country \
4702 east africa|hunter gatherer|indigenous|rift va... English
USA
4958 family relationships|gang|idler|poorhouse|thief English
USA
279                     based on novel English
USA
4244                     based on novel English
USA
4716         china|faith|panama|photography|suspense English
USA
...
...
3355 black comedy|cunnilingus|neo noir|nonlinear ti... English
USA
683 anti establishment|dark humor|fighting|multipl... English
USA
97 ambiguous ending|corporate espionage|dream|sub... English
USA
66 based on comic book|dc comics|psychopath|star ... English
USA
1937 escape from prison|first person narration|pris... English
USA

content_rating
4702 R
4958 R
279 R
4244 R
4716 PG-13
...
...
3355 R
683 R
97 PG-13
66 PG-13
1937 R

[4919 rows x 22 columns]

```

Encode data

For class having dominant groups, replace value to belongs to the dominant group or not For multiple group distributed, do one hot encoding

```
# Replace color column to binary value, true(1) for color and false(0) for black and white
dfFilledData['is_color'] = np.where(dfFilledData['color'] == 'Color', 1, 0)
dfIsColor = dfFilledData.drop('color', axis=1)

# Replace language column to binary value, true(1) for English and false(0) for other
dfIsColor['is_english'] = np.where(dfIsColor['language'] == 'English', 1, 0)
dfIsEnglish = dfIsColor.drop('language', axis=1)

# Replace country column to two binary column, true(1) in is_USA for USA or true(1) in is_UK for UK, other if false(0) for both
countryCat =
dfIsEnglish['country'].where(dfIsEnglish['country'].isin(['USA', 'UK']), 'Other')
countries_dummies = pd.get_dummies(countryCat, dtype=int, prefix='is')
dfSplitCountry : pd.DataFrame = pd.concat([dfIsEnglish,
countries_dummies], axis=1)
dfCountries = dfSplitCountry.drop(columns=['is_Other', 'country'])
```

Split the cell's value as there are some cells contain more than one value Do one-hot encoding and check total numbers of features Built the dataframe after it

```
dfCountries['contentRatingList'] =
dfCountries['content_rating'].str.split(r'\s*\|\s*')
mlbCont = MultiLabelBinarizer()
contEncoded = mlbCont.fit_transform(dfCountries['contentRatingList'])
contEncoded.shape[1]
```

18

```
dfCountries['genresList'] = dfCountries['genres'].str.split(r'\s*\|\s*')
mlbGenre = MultiLabelBinarizer()
genreEncoded = mlbGenre.fit_transform(dfCountries['genresList'])
genreEncoded.shape[1]
```

26

```
dfCountries['keywordsList'] =
dfCountries['plot_keywords'].str.split(r'\s*\|\s*')
mlbKeyword = MultiLabelBinarizer()
keywordEncoded = mlbKeyword.fit_transform(dfCountries['keywordsList'])
keywordEncoded.shape[1]
```

8086

```
# Encoded content_rating and genres; drop keyword as there are more
# than 8000 different value
dfCont = pd.DataFrame(contEncoded, columns=mlbCont.classes_,
index=dfCountries.index)
dfContInData = pd.concat([dfCountries, dfCont], axis=1)

dfGenres = pd.DataFrame(genreEncoded, columns=mlbGenre.classes_,
index=dfContInData.index)
dfGenreInData = pd.concat([dfContInData, dfGenres], axis=1)

dfGenreEncoded =
dfGenreInData.drop(columns=['content_rating','contentRatingList','genres',
'genresList','plot_keywords','keywordsList'])

# Process the target columns, convert from continuous data to class
# data
dfGenreEncoded['imdb_rating'] = pd.cut(dfGenreEncoded['imdb_score'],
bins=[0, 2, 4, 6, 8, 10], labels=[1, 2, 3, 4, 5])
dffeaValue = dfGenreEncoded.drop('imdb_score',axis=1)
```

Handling outliers

Replace outliers with the bounds of the data

```
dfRemoveOutliers = allOutliersToBounds(dffeaValue)
```

Resampling

Constraint data to suitable types

```
# Most of the data columns are counted by people and only one is year
columns = [
    'num_critic_for_reviews',
    'director_facebook_likes',
    'actor_3_facebook_likes',
    'actor_1_facebook_likes',
    'num_voted_users',
    'cast_total_facebook_likes',
    'facenumber_in_poster',
    'num_user_for_reviews',
    'title_year',
    'actor_2_facebook_likes',
    'movie_facebook_likes'
]
# Change the data type to int to prevent data generate by oversampling
# contain float value
dfRemoveOutliersInt = dfRemoveOutliers.astype({col: int for col in
```

```

columns})
dfRemoveOutliersInt

    num_critic_for_reviews  duration  director_facebook_likes \
4702                  1        71.0                      0
4958                  1       110.0                      0
279                   0        55.5                      0
4244                  1       100.0                      0
4716                  0        90.0                      0
...
3355                 ...      155.5                     ...
683                   315      151.0                     459
97                    407      148.0                     459
66                    407      152.0                     459
1937                 199      142.0                     0

    actor_3_facebook_likes  actor_1_facebook_likes  gross \
4702                  21            1000  2.503566e+07
4958                  0              2  3.000000e+06
279                   365              5  2.503566e+07
4244                  338             749  2.503566e+07
4716                  271             595  2.503566e+07
...
3355                 ...        13000  1.079300e+08
683                   637        11000  3.702340e+07
97                    1378        26588  1.150246e+08
66                    1378        23000  1.150246e+08
1937                 461         11000  2.834147e+07

    num_voted_users  cast_total_facebook_likes  facenumber_in_poster \
4702                  5                  1359                      1
4958                  5                  4                      1
279                   6                  5                      0
4244                  6                 1814                      0
4716                  6                 1754                      0
...
3355                ...                  ...                      ...
683                 221859                 16557                      1
97                  221859                 13209                      2
66                  221859                 31937                      0

```

1937	221859	13495	0	
<code>Short \ num_user_for_reviews ... News Reality-TV Romance Sci-Fi</code>				
4702	1	...	0	
0			0	
4958	1	...	0	
0			0	
279	0	...	0	
0			0	
4244	1	...	0	
0			0	
4716	1	...	0	
0			0	
...	
...	
3355	704	...	0	
0			0	
683	704	...	0	
0			0	
97	704	...	0	
0			0	
66	704	...	0	
0			0	
1937	704	...	0	
0			0	
<code>Sport Thriller War Western imdb_rating</code>				
4702	0	0	0	4
4958	0	0	0	3
279	0	0	0	4
4244	0	1	0	4
4716	0	1	0	4
...
3355	0	0	0	5
683	0	0	0	5
97	0	1	0	5
66	0	1	0	5
1937	0	0	0	5
[4919 rows x 64 columns]				

Define strategy for resampling

```
X = dfRemoveOutliersInt.drop('imdb_rating', axis=1)
y = dfRemoveOutliersInt['imdb_rating']
```

Scaling

Use minmax scaler to scale all data except target to range 0 to 1

```
ct = ColumnTransformer([('scale', preprocessing.MinMaxScaler(),  
dfRemoveOutliersInt.columns.drop('imdb_rating'))],  
                           remainder='passthrough',  
                           verbose_feature_names_out=False).fit(dfRemoveOutliersInt)  
ct.set_output(transform='pandas')  
dfProcessedData = ct.transform(dfRemoveOutliersInt)
```

Write data into file

```
dfProcessedData.to_csv(f'../Dataset/dataFrameProcessed', index=False)
```

Resampling

```

import os

import numpy as np
import pandas as pd
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline
from matplotlib import pyplot as plt

from sklearn.ensemble import RandomForestClassifier,
HistGradientBoostingClassifier
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import make_scorer, precision_score,
confusion_matrix, ConfusionMatrixDisplay, get_scorer
from sklearn.model_selection import train_test_split, cross_validate,
StratifiedKFold, GridSearchCV
from sklearn.preprocessing import MinMaxScaler

from sklearn.svm import SVC

os.environ['LOKY_MAX_CPU_COUNT'] = str(os.cpu_count()-1) # To silence
warning : Could not find the number of physical cores

randomState = 42

# Function to calculate weighted specificity
def multiclassSpecificity(yTrue, yPredict):
    cm = confusion_matrix(yTrue, yPredict)

    specificities = []
    for i in range(len(cm)):
        trueNegative = np.sum(cm) - np.sum(cm[i, :]) - np.sum(cm[:, i]) + cm[i, i]
        falsePositive = np.sum(cm[:, i]) - cm[i, i]
        specificity = trueNegative / (trueNegative + falsePositive) if
(trueNegative + falsePositive) > 0 else 0
        specificities.append(specificity)

    return np.mean(specificities)

```

Read data and define target

```

df = pd.read_csv('../Dataset/dataFrameProcessed')
targetCol = 'imdb_rating'

```

Split features and target data

```

X = df.drop(columns=[targetCol])
y = df[targetCol]

```

Built set of models selected

```

models = {
    'Support Vector Machine': SVC(kernel='rbf',
random_state=randomState),
    'Random Forest' : RandomForestClassifier(random_state=randomState),
    'Hist Gradient Boosting' :
HistGradientBoostingClassifier(random_state=randomState),
}

```

Built scoring metrix

```

scoringMetrix = {
    'accuracy' : 'accuracy',
    'precision' : make_scorer(precision_score, average='weighted',
zero_division=1),
    'recall' : 'recall_weighted',
    'f1' : 'f1_weighted',
    'specificity' : make_scorer(multiclassSpecificity)
}

```

Split data (80:20)

```

XTrain, XTest, yTrain, yTest = train_test_split(X, y, test_size=0.2,
random_state=randomState , stratify=y)

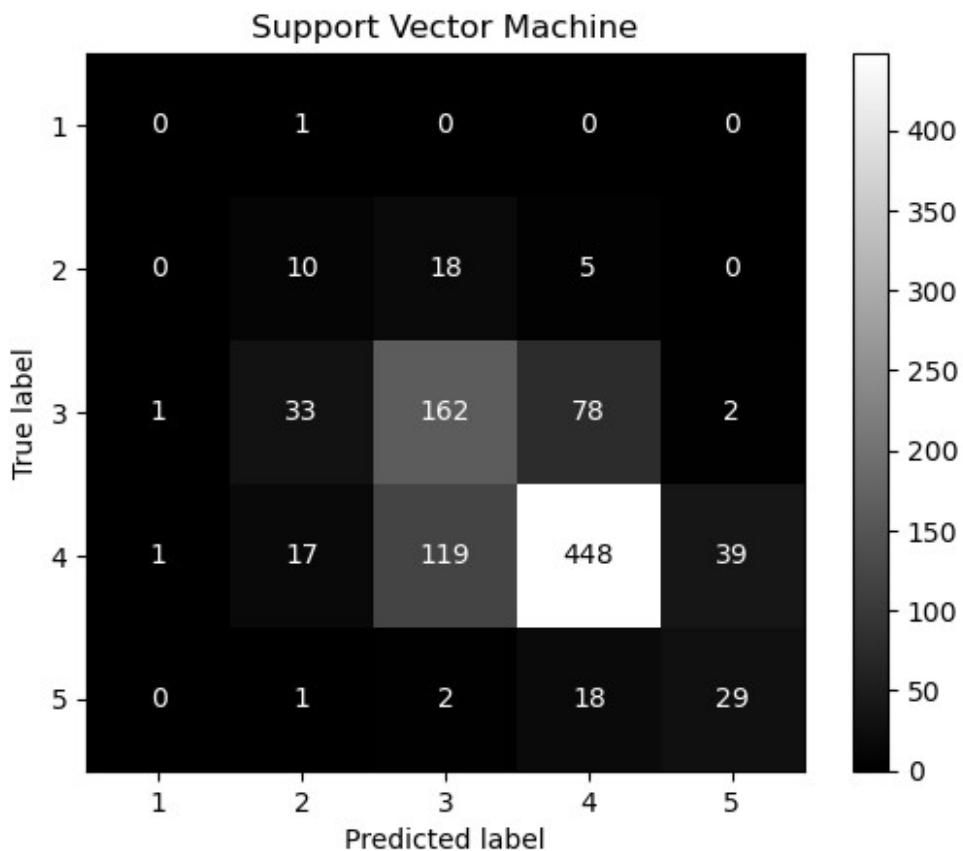
# Function to make pipeline
def makePipeline(modelToUsed):
    steps = [
        ('over', SMOTE(sampling_strategy='auto', random_state=42,
k_neighbors=3)),
        ('scaling', MinMaxScaler().set_output(transform='pandas')),
        ('feature_selection', SelectFromModel(
            estimator=RandomForestClassifier(random_state=randomState,
n_jobs=-1),
            threshold='median')),
        ('classifier', modelToUsed),
    ]
    return Pipeline(steps=steps)

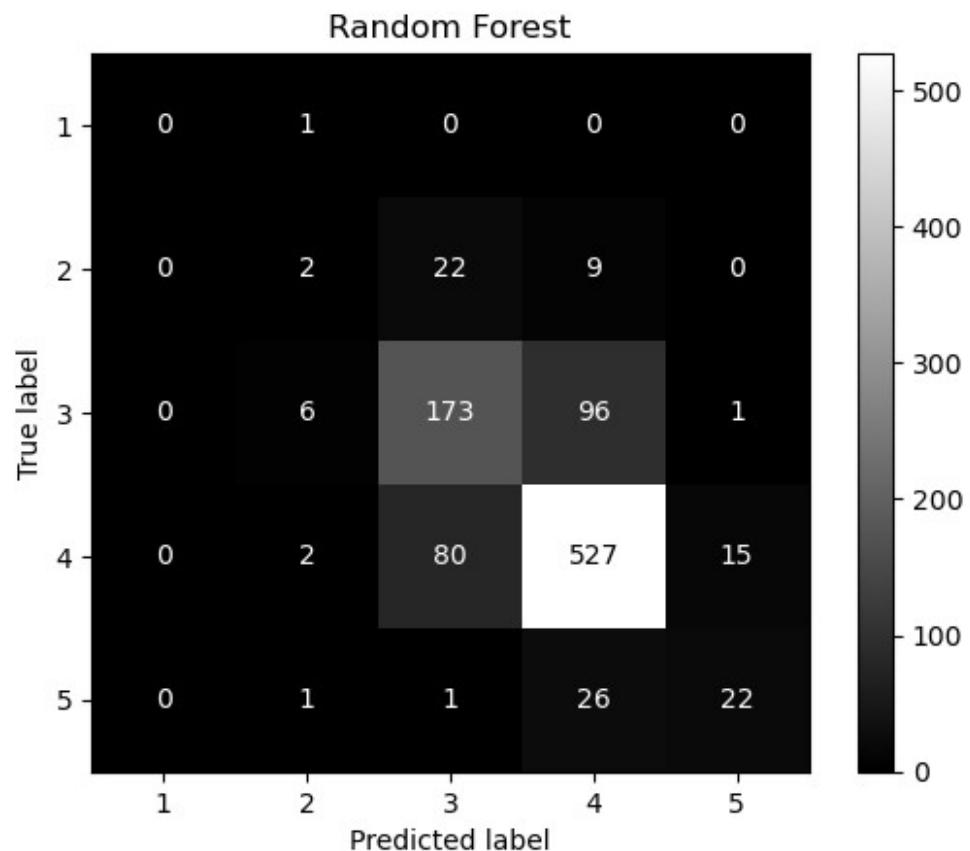
for name, model in models.items():
    # Fit pipeline
    pipeCM = makePipeline(model)
    pipeCM.fit(XTrain, yTrain)

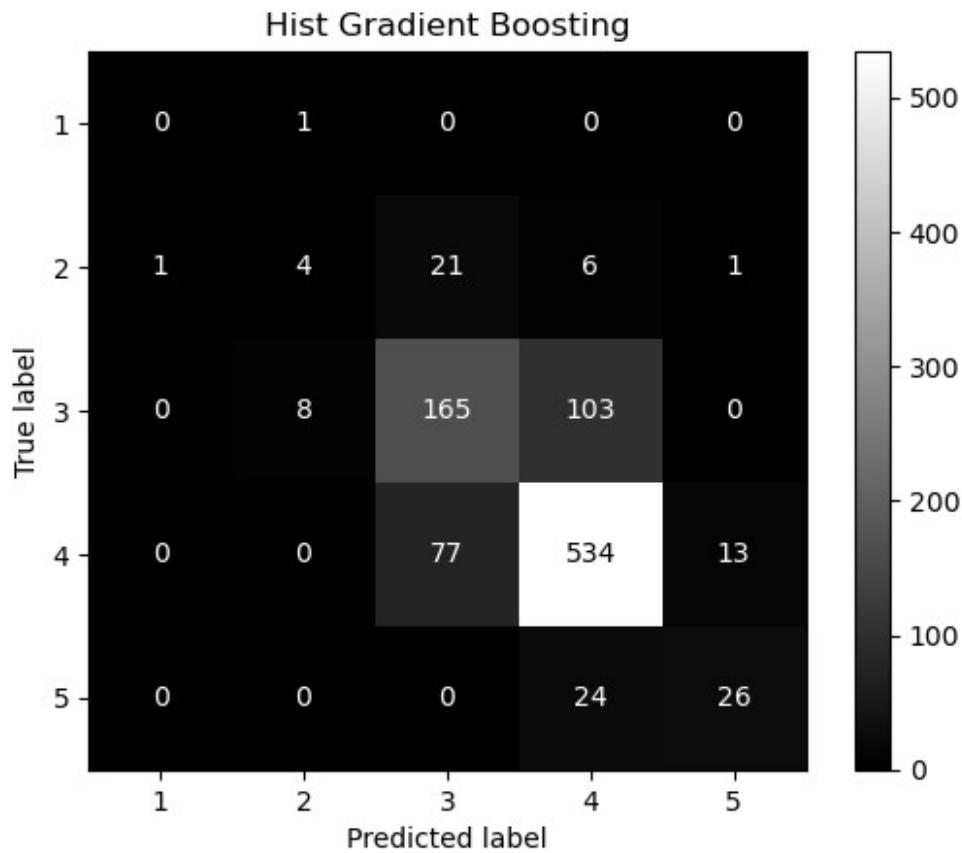
    # Show confusion metrix
    dispCM = ConfusionMatrixDisplay.from_estimator(
        pipeCM,
        XTest,
        yTest,

```

```
cmap='gist_gray',
)
dispCM.ax_.set_title(name)
plt.show()
```







```

for name, model in models.items():
    pipeSC = makePipeline(model)
    pipeSC.fit(XTrain, yTrain)

    # Print model name and its results
    print(f'\n{name}:')
    for metricName, scorer in scoringMetrix.items():
        scorer = get_scorer(scorer)
        score = scorer(pipeSC, XTest, yTest)
        print(f'{metricName}: {score:.5f}')

```

Support Vector Machine

```

accuracy      : 0.65955
precision     : 0.69490
recall        : 0.65955
f1            : 0.67352
specificity   : 0.88450

```

Random Forest

```

accuracy      : 0.73577
precision     : 0.71973
recall        : 0.73577

```

```

f1 : 0.72557
specificity : 0.89260

Hist Gradient Boosting
accuracy : 0.74085
precision : 0.72702
recall : 0.74085
f1 : 0.73153
specificity : 0.89333

dataset_results = []

for name, model in models.items():
    pipeCV = makePipeline(model)

    cvResult = cross_validate(
        pipeCV,
        X,
        y,
        cv= StratifiedKFold(n_splits=5, shuffle=True,
random_state=randomState),
        scoring=scoringMetrix,
        n_jobs=-1,
        error_score='raise'
    )

    dataset_results.append({
        'Model' : name,
        'Accuracy' : np.mean(cvResult['test_accuracy']),
        'Precision' : np.mean(cvResult['test_precision']),
        'Recall' : np.mean(cvResult['test_recall']),
        'F1 Score' : np.mean(cvResult['test_f1']),
        'Specificity' : np.mean(cvResult['test_specificity'])
    })

print(pd.DataFrame(dataset_results).set_index('Model'))

          Accuracy  Precision   Recall   F1 Score
Specificity
Model

Support Vector Machine  0.666801  0.709935  0.666801  0.682750
0.890008
Random Forest         0.730029  0.722563  0.730029  0.724174
0.893109
Hist Gradient Boosting 0.744256  0.734465  0.744256  0.736981
0.896896

pipeGrid =
makePipeline(HistGradientBoostingClassifier(random_state=randomState))

```

```

param_grid = {
    'classifier_learning_rate': [0.1, 0.2],
    'classifier_max_iter': [100, 200],
    'classifier_max_depth': [None, 10],
    'classifier_max_leaf_nodes': [10, 20],
    'classifier_l2_regularization': [0.01, 0.02],
}

gridSearch = GridSearchCV(
    estimator=pipeGrid,
    param_grid=param_grid,
    scoring='f1_weighted',
    cv=5,
    n_jobs=-1,
    verbose=3
)

gridSearch.fit(XTrain, yTrain)

print(gridSearch.best_params_)
for metricName, scorer in scoringMetrix.items():
    scorer = get_scorer(scorer)
    score = scorer(gridSearch, XTest, yTest)
    print(f'{metricName:<12}: {score:.5f}')

yPred = gridSearch.predict(XTest)
confusionMatrix = confusion_matrix(yTest, yPred)
display = ConfusionMatrixDisplay(confusion_matrix=confusionMatrix,
display_labels=gridSearch.classes_)

fig, axes = plt.subplots(figsize=(6, 6))
display.plot(cmap='gist_gray', ax=axes, values_format='d')
plt.title('Tuned Model')
plt.show()

Fitting 5 folds for each of 32 candidates, totalling 160 fits
{'classifier_l2_regularization': 0.02, 'classifier_learning_rate': 0.1, 'classifier_max_depth': 10, 'classifier_max_iter': 200, 'classifier_max_leaf_nodes': 20}
accuracy : 0.74593
precision : 0.73509
recall : 0.74593
f1 : 0.73791
specificity : 0.89556

```

